

Prashanth Koganti

Azure Data Engineer | Microsoft Fabric Specialist

Prashanth.koganti21@gmail.com | 9989253136 | linkedin.com/in/prashanthkoganti | Hyderabad, India

PROFESSIONAL SUMMARY

Azure Data Engineer with 5+ years of experience at NTT Managed Services, delivering end-to-end data platform solutions. Core expertise spans ADF pipeline orchestration, ADLS Gen2 data lake architecture, PySpark transformations in Databricks, and Azure Synapse Analytics. Recent focus on Microsoft Fabric migrations and DirectLake optimisation. Holds AZ-104 and DP-700 certifications. Experienced in Terraform-based infrastructure provisioning, RBAC security, Microsoft Purview data governance, and cross-functional collaboration with BI, security, and client stakeholder teams.

PROFESSIONAL EXPERIENCE

NTT Managed Services Pvt Ltd, Azure Data Engineer

03/2021 to Present | Hyderabad, India

- Built and maintained parameterised ADF pipelines ingesting data from SQL Server, Oracle, REST APIs, and flat files into ADLS Gen2; pipelines ran on schedule or event triggers depending on the source.
- Designed a three-layer ADLS Gen2 data lake (Raw, Curated, Consumption) — managed folder organisation, access policies, and lifecycle rules to control storage spend.
- Developed PySpark notebooks in Databricks for complex transformations: deduplication, column casting, null handling, and reference data joins before loading into the Gold Delta layer.
- Created external tables, stored procedures, and SQL pool optimisations in Azure Synapse Analytics to support BI reporting; resolved partition mismatches, missing statistics, and result set cache settings.
- Migrated Power BI datasets to Microsoft Fabric Lakehouses and switched reports to DirectLake connections for two internal reporting teams, eliminating the need for scheduled imports.
- Implemented Event Hubs and Stream Analytics pipelines for IoT feed processing; output routed to ADLS for historical storage and into SQL tables queried directly by the operations team.
- Provisioned and maintained Terraform modules for ADF, ADLS Gen2, Synapse, and Databricks workspaces across Dev, UAT, and Production environments with separate state files per environment.
- Configured RBAC, AAD groups, private endpoints, and Azure Policies to enforce tagging standards and resource location restrictions across all data services.
- Wrote KQL queries in Azure Monitor and Log Analytics to track pipeline run status, data volumes, and ingestion latency; maintained alert thresholds as workloads evolved.
- Implemented Microsoft Purview — registered and scanned data sources, tagged sensitive columns, and traced lineage across ADF pipelines to support a client GDPR audit.
- Automated recurring admin tasks (clearing stale pipeline runs, checking storage usage, rotating access keys) using Python and PowerShell scripts.
- Participated in sprint planning, code reviews, and documentation; maintained runbooks enabling the on-call team to resolve common pipeline failures without escalation.

TECHNICAL SKILLS

Cloud Platform	Microsoft Azure
Data Integration	Azure Data Factory, Event Hubs, Stream Analytics
Analytics & Warehouse	Azure Synapse Analytics, Microsoft Fabric (Lakehouse, Warehouse, OneLake), Azure Databricks
Storage & Databases	ADLS Gen2, Azure Blob Storage, Azure SQL Database, Cosmos DB, Delta Lake
IaC & Automation	Terraform, Azure CLI, PowerShell, ARM Templates
Programming & Query	PySpark, Python, SQL, T-SQL, Spark SQL, KQL
Governance & Security	Microsoft Purview, Azure Active Directory, RBAC, Azure Policies, Private Endpoints
Monitoring	Azure Monitor, Log Analytics
Reporting	Power BI, Microsoft Fabric Reports, DirectLake, Semantic Models

CERTIFICATIONS

Microsoft Certified: Azure Administrator Associate (AZ-104)
Microsoft Certified: Fabric Analytics Engineer Associate (DP-700)

Microsoft
Microsoft

PROJECTS

Enterprise Data Lake Consolidation

NTT Managed Services

A large manufacturing client needed 12 source systems (SAP, Salesforce, on-premises SQL Server, and flat file feeds) consolidated into a single central data lake.

- Parameterised the ADF pipeline framework with a config table in Azure SQL — onboarding a new source required only a table update, not pipeline rewrites; reduced source onboarding from 2–3 days to a short test run.
- Incremental loads used watermark columns and CDC from SQL Server; sources without CDC support received full weekly loads instead.
- Silver layer transformations in Databricks covered deduplication, type casting, null handling, and reference joins; Gold Delta tables tuned to cut average BI query time from ~8 minutes to under 90 seconds.
- Configured private endpoints for ADLS and ADF (no public internet traffic); Purview scanning applied across all ADLS containers and SQL pools with a custom classification rule for customer ID columns flagged during GDPR audit.

Microsoft Fabric Migration and Reporting Modernisation

NTT Managed Services

A retail client wanted to move off Power BI Premium to reduce licensing costs and consolidate data storage onto Microsoft Fabric.

- Migrated 40+ Power BI datasets into Fabric Lakehouses and reconfigured reports to DirectLake mode, removing the need for scheduled imports.
- Replaced a mixed ingestion layer (Dataflows Gen1 and manual CSV uploads) with ADF pipelines writing directly into OneLake.
- Set up a Fabric Warehouse with T-SQL views and stored procedures for the finance team to handle cross-database joins and replace manual Excel-based monthly calculations.
- Configured Azure Monitor refresh alerts; handed over the full solution to the client's internal team with documentation and a walkthrough session.

KEY ACHIEVEMENTS

- Gold layer average query time reduced from ~8 minutes to under 90 seconds after optimising Delta tables in Databricks.
- Monthly Azure data spend cut by ~22% through Databricks cluster right-sizing and enabling auto-pause on idle Synapse dedicated pools.
- New data source onboarding reduced from 2–3 days to a config table update and a short test run via the parameterised ADF framework.
- Retail client eliminated Power BI Premium capacity costs by migrating to Microsoft Fabric with DirectLake connections.
- Purview implementation gave the client full data lineage visibility and sensitive column coverage across all data assets — the primary requirement of the GDPR audit.

EDUCATION

Sri Indu Institute of Engineering and Technology, B.Tech in ECE
Sri Chaitanya Jr College, Intermediate
Siddarth Model High School, Secondary School Certificate

08/2016 to 2020 | Hyderabad
06/2014 to 04/2016 | Ameenpur
2003 to 03/2014 | Medak

DECLARATION

I hereby declare that the information provided above is true and correct to the best of my knowledge and belief.

Prashanth Koganti
Hyderabad